
NeuroForge: A Self-Correcting, Geometry-Native Neural CFD Engine with Calibrated Physics-Residual Trust

Ali Jabbary

Independent Researcher

research@alijabbary.com

Abstract

Machine-learning surrogates for computational fluid dynamics (CFD) predict steady flow fields over engineering geometries orders of magnitude faster than classical solvers, but they emit a single field with no built-in way for a user to know whether to trust it—especially out of distribution, where early-design exploration lives. We close the loop with the governing physics: we compute the discretised steady-RANS residual of the prediction and ask what jobs it can do. Our central finding is a clean three-way dissociation, validated on a SOTA Transolver backbone: **the physics residual is an excellent, backbone-agnostic trust signal and a useful input feature to a learned corrector, but a poor correction objective.** (1) *As a trust signal*, the residual is a calibrated, backbone-agnostic proxy for error: its rank correlation with field error is 0.625 ± 0.019 on the SOTA backbone (vs ≈ 0.40 on a weak grid backbone, which a learned corrector lifts to 0.83), and a split-conformal layer attains target coverage (0.902 ± 0.008 at the 0.90 target). (2) *As an input feature to a learned, supervised corrector*, the residual helps: a deep-equilibrium (DEQ) corrector that consumes the residual as input and trains toward ground truth reduces volume-field MSE on all three seeds— mse_u -9% , mse_v -21% , mse_p -25% —on a backbone that is itself SOTA-competitive ($\rho_{C_I} = 0.999$, $\approx 5\%$ force error). (3) *As a correction objective or acceptance gate*, the residual fails: sweeping the iterations *raises* the PDE residual while *lowering* field error, and a backtracking test that admits only residual-reducing steps accepts essentially none—minimising the residual is not minimising the error. We package a **distribution-free conformal trust layer** that attains the coverage above (a guarantee that holds only under exchangeability) and empirically retains target coverage out-of-distribution via uncertainty inflation; pairing it with a deep-ensemble σ yields an *input-adaptive* band on the near-deterministic SOTA model (ECE 0.074, $q = 2.35$), where MC-dropout gives a valid but near-constant one. We report one honest caveat plainly: correction quality is *backbone-dependent* (the same corrector is flat on the weak grid backbone). The contribution is the self-auditing trust layer and the residual’s three roles, demonstrated on a competitive surrogate.

Keywords: neural operators, surrogate CFD, physics residuals, uncertainty quantification, conformal prediction, trust calibration, airfoil aerodynamics.

1 Introduction

Classical CFD resolves the governing PDEs by iterating a discretised system until its residuals fall below a tolerance. It is accurate and trusted, but slow: meshing and a converged steady-RANS solve over a new geometry can take minutes to hours, throttling the rapid design-space exploration that early engineering most needs. Neural surrogates promise the opposite trade—millisecond inference—and a vigorous literature now maps geometry plus boundary conditions directly to flow fields with neural operators and attention models.

The problem is **trust**. A trained surrogate is a one-shot function: it emits a field with no self-assessment, no convergence signal, and no recovery path when it extrapolates. On a geometry far from the training set—the common case in design—the user cannot distinguish a faithful prediction from a confident hallucination.

The governing physics offers an obvious lever. The steady incompressible RANS residual of any candidate field is computable directly from the field, with no ground truth. A neural-CFD pipeline could use this residual in three distinct ways: as a *trust signal* that flags where the surrogate is wrong (and drives a conformal trust layer); as an *input feature* to a learned corrector that consumes the residual map and is trained toward ground truth; and as a *correction objective or acceptance gate* that a step must reduce. These are not the same role, and—this is the crux—they do not earn the same verdict. This paper separates them and answers each empirically on a SOTA backbone.

Research question. *In which of its three roles—trust signal, input feature to a learned corrector, and correction objective/gate—is the steady-RANS physics residual of a neural-CFD prediction useful, and does the answer survive on a competitive backbone?*

Thesis (one sentence). *The steady-RANS physics residual is a reliable, conformally-calibrated, backbone-agnostic trust signal and a useful input to a learned corrector that measurably improves a SOTA surrogate—even though minimising the residual directly remains a poor correction objective.*

We arrive at this thesis the hard way. We built the full predict→verify→estimate-uncertainty→correct→fall-back engine, pre-registered hypotheses, and ran a 3-seed ablation on real AirFRANS plus an out-of-distribution study, two formal certificates, and—in a second phase—the same trust layer and corrector on a SOTA Transolver backbone. The trust signal and the supervised corrector survive on the strong backbone; the residual-as-objective machinery (the iteration sweep and the acceptance gate) is the part that disappoints, and we report that as a finding rather than hide it. A key honest finding cuts across all three roles: *correction quality is backbone-dependent*—the supervised corrector that helps on the SOTA backbone is flat on the weak grid backbone—so we report both.

1.1 Contributions

In priority order:

1. **The trust layer is backbone-agnostic (headline).** The physics residual is a strong, calibrated error detector across backbones. On the SOTA Transolver backbone its rank (Spearman) correlation with field error is 0.625 ± 0.019 (subsection 5.2, Table 1); on a much weaker grid backbone it is ≈ 0.40 bare and a learned corrector lifts it to 0.83 (subsection 5.3), so the detector survives both regimes. Building on it, a split-conformal trust layer attains target coverage on the SOTA backbone (0.902 ± 0.008 at the 0.90 target, subsection 5.2); with a deep-ensemble σ the band is input-adaptive on that same near-deterministic model (ECE 0.074, Table 2), and on the weak backbone it holds under distribution shift, including a regime where the bare model’s signal collapses (subsection 5.4, subsection 5.6). The layer depends only on a model’s predictions and uncertainty, so it is backbone-agnostic.
2. **Self-correction works on a strong backbone.** A learned, residual-driven DEQ corrector that consumes the residual as an *input feature* and is trained (supervised) toward ground truth reduces volume-field MSE on the SOTA backbone on **all three seeds**: mse_u $0.127 \rightarrow 0.116$ (−9%), mse_v $0.100 \rightarrow 0.079$ (−21%), mse_p $630 \rightarrow 471$ (−25%), with force ranking held ($\rho_{C_i} = 0.999$) and force error slightly improved (subsection 5.2, Table 1). The backbone is itself SOTA-competitive, so the corrector improves a competitive surrogate.
3. **Residual-as-objective still fails (the surviving negative).** Used as a *correction objective or acceptance gate*—rather than an input feature—the residual does not work. Sweeping the correction iterations *raises* the PDE residual ($0.11 \rightarrow 0.62$) while *lowering* field error (mse_u $3.92 \rightarrow 2.29$): the two move in opposite directions (subsection 5.5, Figure 4). A backtracking acceptance test

that admits only residual-reducing steps accepts essentially zero steps on a trained corrector (subsection 5.5). Minimising the residual is not minimising the error—which is exactly why role (2) supervises toward truth rather than minimising the residual.

4. **Backbone-dependence, reported as a finding (not hidden).** The same corrector that delivers -9 to -25% on the SOTA backbone is flat-to-unhelpful on the weak grid backbone (subsection 5.3, Table 3), where it trades volume accuracy for surface fidelity and a stronger trust signal. Correction quality is therefore backbone-dependent; we report both regimes side by side and the SOTA result is the one that supports the self-correction claim. For honest context we also report a matched-budget baseline study (subsection 5.7, Table 6): the grid backbone trails Transolver by $4\text{--}60\times$, while the trust layer and corrector are demonstrated *on* the competitive Transolver backbone in subsection 5.2.
5. **A reproducible, CPU-first open-source package (neuroforge-cfd)** with a frozen I/O contract, interchangeable backbones, a zero-download synthetic potential-flow data generator, an AirFRANS loader, the ablation/certificate harness, and the figures—so every claim runs on a laptop and every certificate is reproducible from a committed script.

We are explicit throughout about what is **measured** versus **assumed**, retain the retraction of stale single-seed numbers from earlier drafts (subsection 5.3), and make no out-of-distribution coverage *guarantee*; we make a competitiveness claim only for the Transolver backbone our trust layer runs on (subsection 5.2), not for the grid backbone.

2 Related Work

Neural operators (FNO / Geo-FNO). Fourier Neural Operators learn a resolution-agnostic mapping between function spaces by parameterising a global convolution in the spectral domain (Li et al., 2021). Geo-FNO (Li et al., 2022) extends this to irregular geometries via a learned deformation to a latent uniform grid where the FFT is valid. These are strong *one-shot* operators with no self-assessment; NeuroForge implements both as backbones and adds the trust layer around them.

Point-cloud and physics-attention operators (DoMINO, Transolver). DoMINO (Ranade et al., 2025) is a decomposable multi-scale point-cloud operator for external aerodynamics on large industrial meshes. Transolver (Wu et al., 2024) replaces quadratic attention with attention over learnable *physics slices*, giving linear cost and strong accuracy on PDE benchmarks; Transolver++ (Luo et al., 2025) scales it to million-scale meshes. We use Transolver in two ways: as a matched-budget baseline for our grid backbone (subsection 5.7), and as the SOTA backbone our trust layer and corrector actually run on (subsection 5.2), where the residual is a calibrated detector and the supervised corrector reduces field MSE.

Learned solver-correctors and the “fixer” premise. Several lines of work use the discretised residual or a coarse-solution error as a *correction objective*: learned PDE solvers with convergence guarantees (Hsieh et al., 2019), deep-equilibrium operators for steady PDEs (Marwah et al., 2023), iterative refiners (Lippe et al., 2023), and learned residual-error correctors (Anonymous, 2023). PINNs (Raissi et al., 2019) embed the residual in the training loss. NeuroForge implements a contractive DEQ corrector and a feed-forward residual-conditioned corrector squarely in this family. **Our contribution to this line is to separate two roles the prior work conflates.** As a correction *objective*—where a step must reduce the residual—residual minimisation does not track error minimisation on a real RANS benchmark (the residual rises as error falls; the acceptance gate accepts almost nothing). As an *input feature* to a corrector *supervised toward ground truth*, the residual is useful: on a SOTA backbone our DEQ corrector reduces field MSE by $9\text{--}25\%$ (subsection 5.2). The distinction—residual-as-input helps, residual-as-objective does not—sharpens *how* the residual-as-correction premise should be used.

Uncertainty quantification and conformal prediction for PDE surrogates. Deep ensembles (Lakshminarayanan et al., 2017) and MC-dropout (Gal & Ghahramani, 2016) are standard epistemic estimators; conformal prediction gives distribution-free coverage and has been applied to operator learning (Ma et al.,

2024). We do **not** propose a new UQ method. Our trust layer applies split-conformal calibration to MC-dropout uncertainty; the new content is the *characterisation* of when this holds for CFD surrogates—in particular that coverage survives out-of-distribution empirically via uncertainty inflation, while the formal guarantee does not transfer once exchangeability fails.

Benchmarks (AirfRANS). AirfRANS (Bonnet et al., 2022) provides ~ 1000 incompressible steady-RANS simulations over NACA airfoils with **full**, **scarce**, **reynolds**, and **aoa** splits designed to probe generalisation. We use **full** for in-distribution evaluation and the regime-disjoint **reynolds/aoa** splits for the out-of-distribution study.

2.1 Positioning

The closest prior art treats the physics residual primarily as a *correction objective* (PINNs, FNO-DEQ, learned correctors) or treats UQ in isolation from physics. NeuroForge’s distinctive question is which of three roles the residual can play, evaluated head-to-head on the same model and data and then on a SOTA backbone: as a trust *signal* (detect error, drive conformal calibration), as an *input feature* to a supervised corrector, and as a correction *objective/gate*. Our answer—signal yes, input-feature yes (on a strong backbone), objective no—and the resulting calibrated trust layer are the contribution. The engine’s no-harm/contraction machinery (section 3) is retained because it is correct and reproducible; the acceptance-gate guarantee in particular is real but, as we show, vacuous as an accuracy mechanism.

3 Method

3.1 Pipeline overview

The engine wraps a one-shot backbone in a verify–estimate–correct–fall-back loop. Throughout, the physics residual plays **three distinct roles with three distinct verdicts** (section 5): (i) as a *trust signal* it detects error and drives the conformal layer (excellent, backbone-agnostic); (ii) as an *input feature* to a corrector *supervised toward ground truth* it improves a SOTA backbone (subsection 5.2); (iii) as a correction *objective or acceptance gate* it fails (subsection 5.5). We describe all components for completeness and flag, at each one, which role it instantiates.

```
CAD / STL / airfoil + BCs
|
v geometry-native encoding (SDF + solid mask + coords + freestream + log Re)
x in R^{7 x H x W}
|
v neural-operator / transformer backbone f_theta
yhat in R^{4 x H x W} (u, v, p, nu_t) <-----+
|
v physics residual checker R(.) | correction
continuity, momentum_x, momentum_y, BC violation | loop
| | (analysed, limited)
v uncertainty sigma(.) (MC-dropout / deep ensemble) |
|
v conformal trust layer + trust map T |
| (the surviving contribution) |
v correction operator c_phi -> Delta -----+
| (DEQ fixed point, or feed-forward under acceptance test)
v (only if a region stays low-trust)
uncertainty-gated classical CFD patch (interface)
|
v flow field . Cp . forces . uncertainty/residual/trust maps . history
```

3.2 Geometry-native encoding and governing equations

A **FlowCase** (geometry + boundary conditions + fluid + domain) is encoded into a fixed channel-first stack $x \in \mathbb{R}^{7 \times H \times W}$ in the frozen **INPUT_CHANNELS** order (**sdf**, **mask**, x , y , u_{in} , v_{in} , $\log Re$): the signed distance to the body surface (negative inside the solid), a fluid/solid mask, normalised cell coordinates, the freestream

velocity broadcast over the grid, and $\log_{10}$ Re. The backbone outputs `OUTPUT_CHANNELS` (u, v, p, ν_t) , where p is the **kinematic** pressure p/ρ .

The verifier evaluates the steady, incompressible, 2-D RANS equations in primitive form on the **physical (denormalised)** fields, with effective viscosity $\nu_{\text{eff}} = \nu + \nu_t$, pointwise:

$$\text{continuity: } r_c = \frac{\partial u}{\partial x} + \frac{\partial v}{\partial y}, \quad (1)$$

$$x\text{-momentum: } r_x = u \frac{\partial u}{\partial x} + v \frac{\partial u}{\partial y} + \frac{\partial p}{\partial x} - \nu_{\text{eff}} \nabla^2 u, \quad (2)$$

$$y\text{-momentum: } r_y = u \frac{\partial v}{\partial x} + v \frac{\partial v}{\partial y} + \frac{\partial p}{\partial y} - \nu_{\text{eff}} \nabla^2 v. \quad (3)$$

Because p is kinematic, no explicit density appears. Derivatives use central finite differences with one-sided stencils at the borders; residuals are zeroed inside the solid. A boundary-condition violation map r_{bc} adds a **no-slip** term penalising velocity magnitude in the thin fluid band adjacent to the wall (weighted by $\exp(-|\text{sdf}|/\ell)$, $\ell \approx 3$ cells) and a **far-field** term penalising deviation from (u_∞, v_∞) on the outer border ring. We stress the standard caveat that drives our whole study: a low residual is **necessary but not sufficient** for correctness—a smooth near-freestream field can have near-zero residual yet be entirely wrong—so the residual is a *consistency monitor*, and whether it tracks error is the empirical question [section 5](#) answers.

3.3 Trust map and conformal calibration (the surviving contribution)

Raw trust map. The combined PDE-residual magnitude $\rho_{\text{res}} = \sqrt{r_c^2 + r_x^2 + r_y^2 + r_{bc}^2}$ and a per-cell predictive uncertainty σ are each mapped to $[0, 1]$ (using an absolute physical reference scale $s = U_\infty^2/L + U_\infty/L$ when available, else a robust 95th-percentile normalisation with an absolute floor) and fused,

$$e = \text{clip}(w_r \hat{r} + w_u \hat{\sigma}, 0, 1), \quad T = 1 - e, \quad (4)$$

with $w_r = 0.6$, $w_u = 0.4$, giving a traffic-light field (green $e < 0.15$, red $e > 0.45$, else yellow). Uncertainty σ is the channel-mean standard deviation from a deep ensemble or MC-dropout.

Split-conformal calibration. Raw σ is uncalibrated, so a threshold on it carries no guarantee. We add split-conformal calibration (cf. [Ma et al., 2024](#)): on a held-out calibration set we compute per-channel nonconformity scores $s = |\hat{y} - y|/\sigma$ and take the finite-sample-corrected $(1 - \alpha)$ quantile q ; the band $q \cdot \sigma$ then satisfies the distribution-free coverage guarantee in [Proposition 1](#).

Proposition 1 (Distribution-free coverage under exchangeability). *If the calibration scores and a test score are exchangeable, then the split-conformal band $q \cdot \sigma$ satisfies $\mathbb{P}(|\hat{y} - y| \leq q \sigma) \geq 1 - \alpha$.*

This converts the trust threshold from a hand-picked constant into a band with a known in-distribution coverage level. [subsection 5.6](#) reports the measured coverage, the coverage-vs- α sweep, and the out-of-distribution behaviour, with the exchangeability caveat made explicit.

3.4 The correction loop: input-feature corrector vs. objective/gate

The engine offers two correction operators and a backtracking acceptance test, and the distinction between them *is* the intellectual core of the paper. The DEQ corrector uses the residual as an *input feature* and is trained by supervised regression toward ground truth, applied directly without any acceptance test; this is the path that helps on a SOTA backbone ([subsection 5.2](#)). The backtracking acceptance test uses the residual as an *objective/gate*—a step is admitted only if it lowers the residual norm; this is the path that fails ([subsection 5.5](#)). We describe both precisely so the contrast is interpretable.

Feed-forward corrector + backtracking acceptance test. A small residual CNN c_ϕ predicts an additive correction conditioned on the current field, its current 3-channel residual map, and the geometry:

$\Delta_k = c_\phi(y_k, R(y_k), x)$. A candidate $y_{k+1} = y_k + s \Delta_k$ is accepted only if it does not increase the residual norm $N(y) = \sqrt{r_c^2 + r_x^2 + r_y^2}$:

$$N(y_k + s \Delta_k) \leq N(y_k) + \varepsilon, \quad (5)$$

with s halved up to four times on failure. By construction $N(y_0) \geq N(y_1) \geq \dots \geq N(y_K)$ —the residual norm is monotone non-increasing across accepted steps. This is the convergence discipline of a classical solver. Its weakness is exactly the one we exploit as a finding: the test is satisfiable by the identity (zero step), so a corrector that cannot reduce the residual is simply rejected. Empirically (subsection 5.5) a trained corrector accepts essentially no steps, because reducing the field error and reducing the PDE residual are not the same objective.

Contractive DEQ corrector. To give the loop a genuine convergence guarantee in its *own* iteration variable, the correction δ is defined as the fixed point of a learned operator

$$\delta^* = T_\theta(\delta^*; c), \quad T_\theta(\delta; c) = \kappa \cdot g_\theta([\delta, c]), \quad c = [\hat{y}, r(\hat{y}), \text{geom}], \quad (6)$$

where g_θ is a CNN whose layers are spectrally normalised (each ≤ 1 -Lipschitz) and $\kappa < 1$.

Proposition 2 (Banach contraction of the corrector). *With g_θ each-layer 1-Lipschitz and $\kappa < 1$, T_θ is a κ -contraction in δ . By the Banach fixed-point theorem the equilibrium δ^* exists, is unique, and the iteration converges geometrically, $\|\delta_k - \delta^*\| \leq \kappa^k \|\delta_0 - \delta^*\|$.*

The construction underlying Proposition 2 is a Deep-Equilibrium model (Bai et al., 2019) with the Lipschitz constant controlled by spectral normalisation (Winston & Kolter, 2020), trained with Jacobian-Free Back-propagation (Fung et al., 2022). Crucially, g_θ is trained on **data** (it targets the correction toward truth, with the residual as an input *feature*) rather than by minimising the residual; at inference the converged δ^* is applied directly, **without** the backtracking acceptance test. We verify this contraction empirically (subsection 5.6: measured factor $0.78 < \kappa = 0.9$). Two scopes must be kept distinct, and subsection 5.5 leans on the distinction: the DEQ contraction is in the *correction variable* δ , whereas the *PDE residual of the output field* is a different quantity that, as we show, can rise even as δ converges and the field error falls.

The feed-forward corrector is trained (backbone frozen) so that $c_\phi(\hat{y}_{\text{norm}}, R_{\text{norm}}, x) \approx y_{\text{norm}}^* - \hat{y}_{\text{norm}}$, with the final convolution initialised near zero for a stable first iteration.

3.5 Uncertainty-gated classical fallback

After the loop, if the maximum uncertainty exceeds a threshold, the low-trust fluid region is handed to a `ClassicalFallback`. In the present release this is an **interface with a stub backend** that reports what would run (and the region size) without invoking an external solver; the `openfoam/su2` backends raise `NotImplementedError` with setup guidance. This keeps the engine importable and runnable with nothing installed while fixing the integration seam; it is not part of any quantitative claim.

4 Implementation

NeuroForge is an open-source, **CPU-first**, pure-Python package (`neuroforge-cfd`, Python ≥ 3.10 , NumPy/SciPy/PyTorch/Matplotlib). Importing the package caps BLAS/OMP/MKL thread counts to one by default to avoid catastrophic oversubscription on low-core machines, and does no heavy work. The codebase is organised around a **frozen I/O contract** in `core/`.

Module map. `core/` holds the frozen data contracts (`FlowCase`, `FlowField`, `Diagnostics`, `SolveResult`) and `Config`. `geometry/` builds the 7-channel input (NACA generation, SDF/mask rasterisation, `encode_case`). `data/` has the synthetic generator, the `AirFRANS` loader, the rasteriser, and the `datamodule` (`Normalizer/loaders`). `models/` provides `FN02d`, `GeoFN0`, a Transolver-style `PhysicsTransformer`, `UNet/DeepONet` baselines, `LocalCorrectionNet`, the `DEQCorrector`, and the `DeepEnsemble/MCDropoutUQ` wrappers, all in a string-keyed registry. `physics/` has the differential operators, the `PhysicsChecker`, `trust_map`, `force/Cp/error` metrics, the conformal calibration, and the differentiable

`physics_residual_torch`. `solver/` has `Predictor`, `NeuroForgeEngine`, `neural_residual_iteration`, and `ClassicalFallback`. `train/`, `viz/`, `cli.py`, `app/` provide training, plotting/report, the CLI, and a Streamlit UI.

Fixed I/O contract. Inputs are always the 7 channels (`sdf, mask, x, y, uin, vin, log Re`); outputs always the 4 channels (`u, v, p,  $\nu_t$`). Fields are (`ny, nx`) `float32`; network tensors channel-first (`B, C, ny, nx`). Pressure is kinematic; residuals are computed on denormalised fields with $\nu_{\text{eff}} = \nu + \nu_t$. Fixing this contract is what makes the backbones interchangeable and the trust layer backbone-agnostic.

Backbones. `FNO2d` is a faithful spectral FNO; `GeoFNO` adds a geometry-gated conditioning of the lifted features; `PhysicsTransformer` implements Transolver-style physics-slice attention (linear in grid points); `UNet/DeepONet` are baselines.

Training loss. The `CompositeLoss` sums a masked data MSE over fluid cells, a differentiable physics-residual term on the denormalised fields, and a no-slip BC term. The corrector is trained separately with the backbone frozen (subsection 3.4).

Synthetic potential-flow generator (zero-download reproducibility). `SyntheticRANS` superposes a Hess–Smith source-panel potential core (with a Kutta-enforcing bound vortex), an algebraic near-wall no-slip ramp and Gaussian wake deficit, kinematic Bernoulli pressure, and a mixing-length ν_t , with desingularised kernels and light smoothing. The fields are physically plausible and continuity-respecting but **analytic**—a smoke-test/reproducibility substrate, not solver ground truth. **No quantitative claim in this paper rests on synthetic data**; all reported numbers (section 5) are on real AirfRANS.

AirfRANS loader. `load_airfrans` reads each simulation’s point cloud, reconstructs the airfoil loop, rasterises the targets onto a structured crop, and returns (`FlowCase, FlowField`) pairs for all four splits.

5 Experiments

All quantitative results are on **real AirfRANS**. Unless noted, the in-distribution and out-of-distribution ablations use the pre-registered protocol (`benchmarks/ablation.py`, `docs/EXPERIMENTS.md`) over **3 seeds (0, 1, 2)**, reported as mean $\pm$ std (population std, `ddof=0`; sample std at $n = 3$ widens bars by $\approx 1.22\times$). Metrics follow the AirfRANS community protocol (`evaluate_cases`): per-channel volume MSE (lower is better, well-conditioned), surface-pressure MSE on the body, Spearman rank correlation of the force coefficients ρ_{C_l}, ρ_{C_d} (closer to 1 is better—what early-design ranking needs), and **residual_error_spearman** (> 0 means a low residual tracks low error). The ν_t channel is near-degenerate at this scale ($\text{MSE} \approx 5 \times 10^{-8}$ for every arm) and is omitted from the tables (see section 7).

With $n = 3$ we report **per-seed sign-consistency** plus effect-size magnitude rather than p-values, which are meaningless at $n = 3$. An effect called “robust” below is sign-consistent across all 3 seeds with a large effect size and (where stated) non-overlapping per-seed distributions.

5.1 Setup

The backbone is an FNO trained on AirfRANS **full** (800 train / 200 test, 80 epochs). Four arms: **backbone**, **backbone (no physics loss)**, **backbone + local corrector** (feed-forward, with the acceptance test), and **backbone + DEQ corrector**. The certificate runs (subsection 5.6) use a dropout-enabled FNO (width 48, 4 layers, 20 modes, dropout 0.05) with a DEQ corrector ($\kappa = 0.9$, damping 0.5), trained 40 + 15 epochs at resolution 128. The **v2** phase (subsection 5.2) re-runs the trust layer and DEQ corrector on a SOTA Transolver backbone under the same protocol; subsection 5.3–5.6 are the weak-grid-backbone runs that establish the backbone-dependence finding.

Table 1: SOTA Transolver backbone (7.35M params, 80 ep, AirfRANS full), alone vs. with the DEQ self-correction loop. 3 seeds, mean $\pm$ std (population std, ddof=0). Lower MSE is better; ρ closer to 1 is better; resid $\leftrightarrow$ err $\rho > 0$ supports the trust signal; conformal coverage targets 0.90. Source: results/v2/v2_results.json. See subsection 5.2.

metric	mse_u	mse_v	mse_p	C_l rel. err	C_d rel. err	resid $\leftrightarrow$ err ρ
backbone	0.127 $\pm$ 0.009	0.100 $\pm$ 0.004	630 $\pm$ 42	5.6%	7.5%	0.625 $\pm$ 0.019
+ DEQ loop	0.116 $\pm$ 0.003	0.079 $\pm$ 0.005	471 $\pm$ 15	5.2%	6.7%	0.635 $\pm$ 0.021
Δ	-9%	-21%	-25%	-0.4pp	-0.8pp	+0.010

5.2 Backbone-agnostic trust + self-correction on a SOTA backbone (v2)

We re-run the trust layer and the DEQ corrector on a **SOTA Transolver backbone** (7.35M parameters, 80 epochs, AirfRANS full, 3 seeds; corrector trained 20 epochs with the backbone frozen), the strongest model in our study. The protocol and scoring are identical to subsection 5.3; the source is results/v2/v2_results.json. Table 1 reports the backbone alone and with the DEQ self-correction loop.

Self-correction works on a strong backbone. The DEQ corrector reduces volume field MSE on **all three seeds**: mse_u -9% (0.127 $\rightarrow$ 0.116), mse_v -21% (0.100 $\rightarrow$ 0.079), mse_p -25% (630 $\rightarrow$ 471). The per-seed reductions are sign-consistent down on all three seeds for every volume channel. Force ranking is held at $\rho_{C_l} = 0.999$, $\rho_{C_d} 0.995 \rightarrow 0.996$, and the relative force errors improve (C_l 5.6% $\rightarrow$ 5.2%, C_d 7.5% $\rightarrow$ 6.7%). This is the corrector used as an *input-feature path* (supervised toward ground truth, applied without the acceptance test, subsection 3.4). The backbone is itself SOTA-competitive ($\rho_{C_l} = 0.999$, $\approx$ 5% force error, matching Table 6), so the trust layer and corrector are demonstrated *on* a competitive surrogate—not only on the weak grid backbone.

The trust signal is backbone-agnostic. On this SOTA backbone the bare residual-error correlation is already strong (0.625 $\pm$ 0.019, vs $\approx$ 0.40 for the bare weak backbone in Table 3); the corrector barely moves it (0.635 $\pm$ 0.021), because the strong backbone’s residual is already a good detector without correction. The lift-by-corrector effect we report in subsection 5.3 is a property of the *weak* backbone; on the strong backbone the signal is strong from the start. Either way the residual is a reliable detector, which is what “backbone-agnostic” means.

Conformal coverage holds under both uncertainty estimators. Split-conformal calibration of the pressure channel at $\alpha = 0.1$ attains target coverage on this SOTA backbone with *either* uncertainty estimator: 0.902 $\pm$ 0.008 with MC-dropout (3 seeds; results/v2/v2_results.json) and 0.915 with a deep ensemble (single run; results/uq_ensemble/uq_results.json), both at the 0.90 target. Coverage was never the difficulty—**adaptivity** was. With MC-dropout this near-deterministic model has near-degenerate σ at dropout 0.05, so coverage is delivered by a **near-constant** band: the conformal quantile q is enormous ($\approx 1.5\text{--}1.7 \times 10^9$) and the reliability ECE is $\approx$ 0.31.

A deep ensemble recovers an adaptive band on the SOTA model. Replacing MC-dropout with a deep ensemble of $M=5$ independently-trained Transolver members—using the per-cell standard deviation across members as σ —converts the near-constant band into a genuinely input-adaptive one while preserving coverage (Table 2). On the same 100/100 calibration/test split, the conformal quantile drops from $\approx 10^9$ to $q = 2.35$, the reliability ECE improves from $\approx$ 0.31 to **0.074**, and coverage holds at 0.915 (target 0.90, consistent with the $\approx \pm 0.03$ binomial spread at $n_{\text{test}} = 100$). This is a single ensemble run, not a 3-seed average. The ensemble mean is also more accurate than any single member ($\rho_{C_l} = 0.9995$, $\rho_{C_d} = 0.998$, C_l/C_d relative force error 4.4%/5.7%; results/uq_ensemble/uq_results.json). The practical recipe follows: MC-dropout suffices when a constant-width certificate is acceptable, and a deep ensemble is what makes the band track the error on a near-deterministic SOTA surrogate.

Table 2: Conformal calibration of the pressure channel on the SOTA Transolver backbone ($\alpha = 0.1$, 100/100 cal/test split): MC-dropout vs. a deep ensemble ($M=5$). Both attain target coverage; only the deep ensemble yields an input-adaptive band (finite q , low ECE). MC-dropout: 3 seeds (mean $\pm$ std); deep ensemble: single run. Sources: `results/v2/v2_results.json`, `results/uq_ensemble/uq_results.json`.

σ estimator	coverage (target 0.90)	q	ECE ($\downarrow$)	band
MC-dropout ($p=0.05$)	0.902 $\pm$ 0.008	$\approx 1.6 \times 10^9$	≈ 0.31	near-constant
deep ensemble ($M=5$)	0.915	2.35	0.074	adaptive

Table 3: In-distribution AirFRANS full ablation, 3 seeds, mean $\pm$ std (population std). Lower MSE is better; ρ closer to 1 is better; resid $\leftrightarrow$ err $\rho > 0$ supports the trust signal. Bold marks the best arm per column among the physics-loss arms; the physics-loss-free arm (starred where best overall) is reported separately because it removes the physics objective entirely. See Figure 1.

arm	mse_u	mse_v	mse_p	surf. mse_p	ρ_{C_l}	ρ_{C_d}	resid $\leftrightarrow$ err ρ
backbone	3.479 $\pm$ 0.105	0.385$\pm$0.012	2444.8$\pm$120.7	548823 $\pm$ 27946	0.9868 $\pm$ 0.0005	0.895 $\pm$ 0.013	0.397 $\pm$ 0.003
backbone (no phys. loss)	1.995$\pm$0.042	0.323 $\pm$ 0.008*	1963.5 $\pm$ 55.5*	1123649 $\pm$ 104822	0.9920$\pm$0.0002	0.945$\pm$0.008	0.605 $\pm$ 0.016
backbone + local corr.	3.482 $\pm$ 0.058	0.398 $\pm$ 0.007	2469.4 $\pm$ 71.0	541534 $\pm$ 20453	0.9854 $\pm$ 0.0017	0.888 $\pm$ 0.012	0.389 $\pm$ 0.031
backbone + DEQ corr.	3.457 $\pm$ 0.811	0.880 $\pm$ 0.064	3832.7 $\pm$ 331.3	361681$\pm$12273	0.9856 $\pm$ 0.0017	0.923 $\pm$ 0.014	0.827$\pm$0.002

5.3 The weak grid backbone: corrector trades accuracy, lifts the trust signal

This ablation uses the **weak grid backbone**; it establishes the backbone-dependence we report as a finding, and should be read against the SOTA-backbone result in subsection 5.2, where the same corrector *does* reduce field MSE.

On the weak backbone the corrector trades accuracy. The DEQ corrector is **flat** on volume `mse_u` (3.457 vs backbone 3.479; the per-seed deltas are sign-inconsistent, 2/3, so this is noise, not an improvement) and **worse** on `mse_v` (0.385 $\rightarrow$ 0.880, +129%, 3/3, non-overlapping distributions) and `mse_p` (2445 $\rightarrow$ 3833, +57%, 3/3). The feed-forward local corrector helps on **nothing**: it is at-or-worse than the backbone on every volume MSE channel and 3/3 worse on ρ_{C_d} . On this weak backbone the residual-conditioned correctors do not improve volume accuracy; the DEQ arm buys surface-pressure MSE (−34%, 3/3) and a slightly better drag ranking (ρ_{C_d} 0.895 $\rightarrow$ 0.923, 3/3) at a measured cost to volume velocity and pressure. We do **not** claim the corrector improves accuracy in aggregate *on this backbone*—and this is precisely the contrast that makes correction quality backbone-dependent.

The detector succeeds. The same DEQ arm roughly **doubles** the residual-error Spearman correlation, 0.397 $\rightarrow$ 0.827 (3/3, std $\approx$ 0.002, Cohen’s $d \approx$ 143). A low residual tracks low field error far more reliably with the corrector engaged. This is the detector half of the thesis, and it is the largest, cleanest effect in the table.

A control that sharpens the thesis. Removing the physics loss entirely (backbone (no physics loss)) **improves** `mse_u`, `mse_v`, `mse_p`, ρ_{C_l} , ρ_{C_d} , and the residual-error correlation (all 3/3), at the cost of **doubling** surface-pressure MSE (549k $\rightarrow$ 1124k, 3/3). The physics-loss-free backbone is in fact the **best force-ranking model in the table** (ρ_{C_d} 0.945, ρ_{C_l} 0.992)—better than any corrected model, with no correction loop at all. We therefore explicitly do **not** claim the loop delivers best-in-class force ranking. The physics objective (in the loss or the corrector) buys surface-pressure fidelity and a strong trust signal, not volume accuracy or force ranking *on this weak backbone*—the trust-signal role succeeds while the correction role does not, which is exactly the backbone-dependence the SOTA result in subsection 5.2 resolves.

Retraction (retained from prior drafts). An earlier single-seed run reported ρ_{C_l} rising 0.924 $\rightarrow$ 0.958 and surface MSE −25% under the corrector. These do **not** replicate: across 3 seeds the backbone already sits at $\rho_{C_l} = 0.987$ and the DEQ corrector is 0.986 (flat-to-slightly-worse, sign-inconsistent). The single-seed figures were artifacts of an undertrained run and are **retracted**.

Table 4: Out-of-distribution AirfRANS ablation (regime-disjoint train→test), 3 seeds, mean $\pm$ std. The full four-arm table is in `results/full_research/ood/ablation_ood.md`. See Figure 2.

task	arm	mse_u	mse_v	mse_p	surf. mse_p	ρ_{C_l}	ρ_{C_d}	resid↔err ρ
reynolds	backbone	16.218 $\pm$ 0.746	1.187 $\pm$ 0.149	7220.7 $\pm$ 612.2	964641 $\pm$ 61419	0.950 $\pm$ 0.008	0.893 $\pm$ 0.009	0.748 $\pm$ 0.077
reynolds	+ DEQ	5.300 $\pm$ 1.544	1.208 $\pm$ 0.255	6469.5 $\pm$ 438.0	652281 $\pm$ 45390	0.943 $\pm$ 0.008	0.915 $\pm$ 0.009	0.742 $\pm$ 0.041
aoa	backbone	4.312 $\pm$ 0.077	1.310 $\pm$ 0.039	6042.8 $\pm$ 242.5	2000592 $\pm$ 55092	0.963 $\pm$ 0.003	0.926 $\pm$ 0.002	0.314 $\pm$ 0.064
aoa	+ DEQ	3.538 $\pm$ 0.681	1.407 $\pm$ 0.089	4620.6 $\pm$ 170.2	1042711 $\pm$ 78298	0.966 $\pm$ 0.006	0.905 $\pm$ 0.013	0.746 $\pm$ 0.027

Table 5: Iteration sensitivity sweep: `results/sensitivity/iters.csv`. Single sweep (one checkpoint, not seeded). See Figure 4.

n_iters	mse_u	surf. mse_p	residual_norm	resid↔err ρ
0	3.924	541204	0.113	0.423
1	2.460	428560	0.336	0.644
3	2.287	336718	0.542	0.647
5	2.439	308381	0.594	0.675
10	2.570	300298	0.618	0.703
15	2.575	300664	0.620	0.710

5.4 Residuals detect under distribution shift

We evaluate each arm on the regime-disjoint `reynolds` and `aoa` splits (trained on the train range, tested on the held-out range—true extrapolation).

The trust signal is regime-invariant—the single strongest result. On the `aoa` split the bare backbone’s residual–error correlation **collapses to** 0.314, while the DEQ-corrected model holds it at 0.746 (0.314 $\rightarrow$ 0.746, 3/3, Cohen’s $d \approx 8.9$, non-overlapping: every corrected seed ≥ 0.72 , every backbone seed ≤ 0.38). On `reynolds` both are high (≈ 0.74 –0.75). The corrected model thus keeps a regime-invariant ≈ 0.74 trust signal across both shifts, precisely where the one-shot model loses it. The residual is a reliable detector exactly where detection matters most.

Scoping the rest, honestly (protected negatives). The OOD ablation also shows the DEQ corrector reducing the in-distribution→OOD *gap* on volume velocity, volume pressure, and surface pressure (e.g. `reynolds` mse_u 16.2 $\rightarrow$ 5.3, `aoa` surf. mse_p 2.00M $\rightarrow$ 1.04M, both 3/3). We report this as an observation about the loop’s behaviour, **not** as a competitiveness or generalisation guarantee, and we explicitly preserve two negatives. (i) **mse_v is an attenuated-regression artifact, not a gain:** the DEQ corrector’s *absolute* OOD mse_v is worse than the backbone in every regime (0.880/1.208/1.407 vs 0.385/1.187/1.310); its in-dist→OOD gap only looks smaller because its in-distribution mse_v is already inflated. (ii) **Drag ranking is task-dependent:** ρ_{C_d} improves on `reynolds` (0.893 $\rightarrow$ 0.915, 3/3) but **regresses on aoa** (0.926 $\rightarrow$ 0.905, worse on all 3 seeds). We do not claim the loop uniformly preserves force ranking under shift. Methodological caveat: the in-distribution reference is the `full` split (a different training range), so the gap mixes a train-set change with regime shift, and the `aoa` mse_u effect, while 3/3 directional, is seed-0-dominated.

5.5 Residual-as-objective fails: the residual and the error move apart

The most direct evidence that the residual is a poor *objective* comes from sweeping the number of correction iterations on a single trained checkpoint (Table 5).

As the loop iterates, the **PDE residual norm rises monotonically** (0.11 $\rightarrow$ 0.62) while the **field error falls** (mse_u 3.92 $\rightarrow$ 2.29 by iter 3)—the two objectives move in *opposite* directions. This is the clean statement of detector $\neq$ fixer: the very quantity one would minimise to “fix” the field grows while the field gets better. (Note this does not contradict the DEQ contraction of subsection 5.6: the contraction is in the correction variable δ , which converges; the PDE residual of the *output field* is a different quantity, and it is the one that rises.)

The acceptance test accepts almost nothing. On the feed-forward corrector the backtracking acceptance test—which admits only residual-reducing steps—is the honest “certified” version of the loop. Because the test is satisfiable by the identity (zero step), a corrector that cannot reduce the *residual* is simply rejected, and the feed-forward corrector—flat-to-worse on every accuracy metric in Table 3—has essentially no residual-reducing step to offer, so the certified loop makes almost no accepted progress. The corrector that *does* help on the design metrics (DEQ) is precisely the one applied **without** the acceptance test (subsection 3.4). The “certified self-correction” guarantee is therefore real but vacuous as an accuracy mechanism: it correctly refuses to make the residual worse, and in doing so does almost nothing. (This structural argument rests on the acceptance test’s definition plus the Table 3 feed-forward result; we do not have a committed artifact that quantifies the accepted-step count, and flag it as the thinnest leg of the residual-as-objective negative.) We additionally note that for the DEQ corrector the trust-gating and acceptance-test toggles are structural no-ops—the DEQ branch bypasses both—so the only live correction knob is the applied-delta step size (results/sensitivity/toggles.json).

5.6 The conformal trust layer: coverage and contraction

Contraction (H5). Iterating the learned DEQ operator $\delta_{k+1} = T_\theta(\delta_k)$ from a random initialisation on 24 real AirfRANS cases, the measured per-step ratio $\|\delta_{k+1} - \delta^*\|/\|\delta_k - \delta^*\|$ (in the geometric regime, before the $\sim 10^{-7}$ solve floor) has **median 0.78** and maximum 0.86—strictly below the design bound $\kappa = 0.9$ and the falsification threshold 1—and reaches a relative distance of 10^{-5} to the fixed point in a median of 37 steps (results/certificates/h5_contraction.json). The corrector contracts as designed; this is a property of the *operator*, separate from whether minimising the output’s PDE residual helps the field (it does not, subsection 5.5).

In-distribution coverage (H4). With per-cell MC-dropout σ (16 passes) and a 100/100 calibration/test split of the AirfRANS test set, split-conformal calibration at $\alpha = 0.1$ attains per-channel coverage of **0.911** (*u*), **0.928** (*v*), **0.942** (*p*)—all in the $[0.85, 0.95]$ band and conservatively above the 0.90 target, as the $\geq 1 - \alpha$ guarantee requires. The conformal multipliers ($q = 0.77, 1.30, 1.75$) both tighten an over-dispersed and widen an under-dispersed raw σ , so the calibration is doing real work. Coverage tracks $1 - \alpha$ across a sweep of α (Figure 6). The reliability diagram is exact at the target level but over-covers at lower nominal levels (ECE $u/v/p = 0.064/0.093/0.130$), the signature of a heavy-tailed $|\text{error}|/\sigma$ ratio; we therefore claim coverage at the chosen α , **not** full distributional calibration (results/certificates/h4_coverage.json).

The certificate is calibrated on the deployed, corrected field—not the raw backbone. NeuroForge ships the DEQ-corrected field, so the coverage guarantee must hold for *that* output; split-conformal calibration, however, applies to whatever predictor it scores, and the raw-backbone certificate does **not** transfer unchanged through the corrector. We verify this on the dropout-FNO backbone whose DEQ corrector materially changes the field (mean $|\text{corrected} - \text{raw}| \approx 5$), through the *same cases_to_error_sigma* path as H4 but with a corrected predict function (scripts/verify_conformal_corrected_field.py; 15/15 calibration/test). The clean isolation is a *within-run frozen- q contrast*—same fitted q , same input-conditional σ , same cases, only the raw vs. corrected error map—and it shows the direction of the coverage change tracking the direction of the error change: where the corrector *enlarges* the error the raw band drops below the 0.90 target (v 0.896 $\rightarrow$ 0.819, p 0.864 $\rightarrow$ 0.777), and where it *shrinks* the error the raw band over-covers and stays conservative (u 0.856 $\rightarrow$ 0.888). This is the coverage analogue of detector $\neq$ fixer (subsection 5.5): the DEQ operator contracts in its *own* variable δ (H5, median ratio 0.78), but the Banach bound governs the correction, not $|\text{error}|/\sigma$, so it cannot imply coverage survival—and on v/p residual contraction coincides with error *growth*, which is exactly what breaks the raw band. Re-fitting q on the corrected calibration field—at no change to σ , which is input-conditional and identical for both fields—restores coverage toward target where it had collapsed (v 0.819 $\rightarrow$ 0.894, p 0.777 $\rightarrow$ 0.834) and sharpens it where it was loose (u 0.888 $\rightarrow$ 0.874). We therefore calibrate the shipped certificate on the corrected output; split-conformal validity requires only exchangeable corrected-field scores, not that σ be the corrected field’s own uncertainty. These estimates are directional at $n = 15$ (per-case coverage std ≈ 0.15 –0.23, so $\approx \pm 0.05$ standard error), but the same pattern—raw- q under-covers v/p on the corrected field, refitting restores it—reproduces on an independent 20/20 split (results/certificates/probe_conformal_after_deq.json); the residual p under-coverage is

Table 6: Matched-budget baseline on AirfRANS full ($n_{\text{train}} = 800$, 80 epochs, identical rasterisation and `evaluate_cases` scoring), 3 seeds. See Figure 3.

model	n_params	mse_u	mse_v	mse_p	surf. mse_p	ρ_{C_l}	ρ_{C_d}
Transolver (baseline)	7.35M	0.120 $\pm$ 0.005	0.088 $\pm$ 0.013	628.5 $\pm$ 29	9110 $\pm$ 500	0.9992 $\pm$ 0.0002	0.9963 $\pm$ 0.0012
our backbone (Table 3)	—	3.479	0.385	2444.8	548823	0.987	0.895
our backbone + DEQ	—	3.457	0.880	3832.7	361681	0.986	0.923

the small-sample heavy-tail effect already noted at $n = 100$. On the SOTA Transolver backbone the DEQ corrector reduces field error on all channels (subsection 5.2, Table 1), so its raw-field band is *expected* to remain conservative (over-covering) on the deployed v2 field; recalibration matters precisely for correctors that trade field error for residual reduction, as the dropout-FNO does.

Out-of-distribution coverage (empirical, not a guarantee). Calibrating q on the in-distribution full split and evaluating on the OOD splits, coverage stays in the target band: u 0.90/0.91/0.91, v 0.91/0.94/0.92, p 0.92/0.94/0.90 (full / reynolds / aoa) vs the 0.90 target (results/sensitivity/ood_coverage.json, Figure 5). **This is an empirical observation, not a guarantee.** The conformal guarantee assumes exchangeable calibration and test draws, which fails under shift; coverage nonetheless holds because the MC-dropout σ inflates appropriately on shifted inputs, so the same q still brackets the (larger) errors. We report this as a desirable empirical property and explicitly **do not** assert distribution-free coverage out-of-distribution.

The trust layer depends only on a model’s predictions and σ , so it is backbone-agnostic; subsection 5.2 demonstrates it on a SOTA Transolver backbone, where coverage holds at 0.902 $\pm$ 0.008 (target 0.90) and a deep-ensemble σ makes the band input-adaptive on that near-deterministic model ($q = 2.35$, ECE 0.074; Table 2).

5.7 Honest baseline: the grid backbone trails SOTA

Transolver, a SOTA point-cloud physics-attention transformer, trained on the same data with the same budget and scored identically, is far ahead of our grid backbone: relative to the bare backbone, roughly $29\times$ **on mse_u**, $4\times$ **on mse_v** ($\approx 10\times$ against the volume-regressed DEQ arm), $4\times$ **on mse_p**, and $\sim 60\times$ **on surface pressure** ($\approx 40\times$ against the DEQ arm, whose surface MSE is lower), with near-perfect force ranking. Our grid backbone is not a competitive surrogate; we make **no competitiveness claim** for it. The decisive next step follows in subsection 5.2: we put the trust layer and the DEQ corrector *on* this Transolver backbone, where the residual is a calibrated detector ($\rho = 0.625$) and the corrector reduces field MSE by 9–25%. The large grid-backbone gap is itself informative: even a much weaker backbone yields a residual that is a calibrated, shift-robust detector of its own errors.

Note on the two Transolver runs. The Transolver numbers here (Table 6: mse_u 0.120 $\pm$ 0.005, mse_v 0.088 $\pm$ 0.013, mse_p 628.5 $\pm$ 29) and the v2 backbone-alone numbers (Table 1: 0.127 $\pm$ 0.009, 0.100 $\pm$ 0.004, 630 $\pm$ 42) are the *same Transolver architecture at the same budget* (7.35M params, 80 epochs), trained in two separate runs; the differences are within the reported per-seed spread. The v2 run is the one we use as the NeuroForge backbone in subsection 5.2.

5.8 Figures

5.9 Reading against the pre-registered hypotheses

- **H1 (the corrector improves accuracy).** *Backbone-dependent.* On the **weak grid backbone** it is *not supported as pre-registered*: the feed-forward corrector helps on nothing and the DEQ corrector regresses volume mse_v/mse_p (3/3) while flat on mse_u. On the **SOTA Transolver backbone** it is *supported*: the DEQ corrector reduces all three volume MSE channels on all 3 seeds (mse_u -9% , mse_v -21% , mse_p -25% ; subsection 5.2, Table 1) with force ranking held. Correction quality




Figure 1: In-distribution ablation (Table 3): the DEQ corrector lifts the trust signal ($0.40 \rightarrow 0.83$) and surface fidelity while regressing volume `mse_v/mse_p`.




Figure 2: OOD gap and the regime-invariant trust signal (`aoa` $0.31 \rightarrow 0.75$); the bare backbone’s signal collapses, the corrected model’s holds.

is therefore backbone-dependent—a finding, not a hidden negative—and the supervised corrector works where the backbone is strong.

- **H2 (the residual is a valid trust signal).** *Supported, robustly, in- and out-of-distribution.* $\text{resid} \leftrightarrow \text{err}$ $\rho > 0$ for every arm/seed/split; DEQ lifts it to 0.83 in-dist and makes it regime-invariant (≈ 0.74) OOD, rescuing the `aoa` collapse $0.31 \rightarrow 0.75$.
- **H3 (DEQ $\geq$ feed-forward corrector).** *Holds on the design axis only.* DEQ beats the feed-forward corrector on surface MSE, ρ_{C_d} , and the trust signal (3/3), but is worse on volume `mse_v/mse_p`; the two trade off.
- **H4 (conformal coverage).** *Supported in-distribution* (0.91/0.93/0.94 at $\alpha = 0.1$; SOTA backbone 0.902, deep-ensemble band adaptive at ECE 0.074, Table 2); *empirical only OOD* (subsection 5.6).
- **H5 (DEQ contraction).** *Supported* (measured factor $0.78 < \kappa = 0.9 < 1$).

Every number comes from the committed, reproducible harness (`benchmarks/ablation.py`, `scripts/run_certificates.py`, `results/sensitivity/`).

6 The residual floor: a self-consistent monitor is a good detector but a bad fixer

The iteration sweep of subsection 5.5 (Table 5) shows a counter-intuitive fact: driving the monitored physics residual *down* drives the field error *up*. We now explain this exactly. The mechanism is not ill-conditioning per se; it is that the *monitored* residual is self-consistent at a field that is not the truth—so its minimiser sits away from u^* by an irreducible amount. The same decomposition that forces this also explains why the residual is nonetheless an excellent error *detector*: one equation, both verdicts.




Figure 3: Honest matched-budget baseline; the grid backbone trails Transolver ~ 4 – $60\times$ across channels (no competitiveness claim).




Figure 4: The money figure for the headline: PDE residual *rises* while field error *falls* across correction iterations (detector $\neq$ fixer).

6.1 Setup: the monitored operator

Let R_h be the residual operator the engine actually monitors and (in the objective/gate role) minimises: the *discrete* steady incompressible RANS **continuity and momentum** residual on the (n_y, n_x) grid, evaluated on physical fields with $\nu_{\text{eff}} = \nu + \nu_t$. By construction (`Diagnostics.residual_norm`; section 3) it **omits the no-slip wall term**. Write the objective $J(u) = \frac{1}{2}\|R_h(u)\|^2$. Let u^* be the discrete ground-truth field, define the *residual floor* $r^* := R_h(u^*)$, and for any prediction $\hat{u}$ write the error $e := \hat{u} - u^*$ with the first-order expansion

$$R_h(\hat{u}) = r^* + L e + o(\|e\|), \quad L := DR_h(u^*). \quad (7)$$

Theorem 1 (Residual floor: misplaced minimum, quantified floor, and local divergence). *Assume (H1) R_h is the monitored discrete operator above (continuity+momentum, no-slip omitted); (H2) the floor is nonzero, $r^* := R_h(u^*) \neq 0$; and (H3) $L := DR_h(u^*)$ is injective on the admissible subspace. Then:*

- (i) (Misplaced minimum, exact.) *The uniform-freestream field u_∞ satisfies $R_h(u_\infty) = 0$ exactly, so $\min_u J = 0$; whereas $J(u^*) = \frac{1}{2}\|r^*\|^2 > 0$. Hence u^* is **not** a minimiser of J .*
- (ii) (Misplaced stationarity.) *$\nabla J(u^*) = L^\top r^*$, so u^* is a stationary point of J iff $L^\top r^* = 0$; if $r^* \notin \ker L^\top$, gradient flow on J leaves u^* .*
- (iii) (Quantified floor.) *The minimiser of the linearised objective $\frac{1}{2}\|r^* + L e\|^2$ is $e_\infty = -L^+ r^*$, leaving irreducible field error $\|e_\infty\| = \|L^+ r^*\|$, which grows like $1/\sigma_{\min}(L)$ as L becomes ill-conditioned.*




Figure 5: Conformal coverage stays in-band on OOD splits (empirical, via σ inflation).




Figure 6: In-distribution reliability / coverage-vs- α and the measured DEQ contraction ($0.78 < 1$).

- (iv) (Local divergence direction.) *Under residual-gradient flow $\dot{e} = -\nabla J(\hat{u})$, the field error obeys $\frac{d}{dt} \frac{1}{2} \|e\|^2|_0 = -\|Le\|^2 - (Le) \cdot r^*$. There is an open cone of error directions—those with $(Le) \cdot r^* < -\|Le\|^2$ —on which a residual-reducing step strictly **increases** the field error.*

Proof. (i) For a spatially constant field every finite-difference derivative vanishes, so the discrete continuity and momentum residuals are identically zero: $R_h(u_\infty) = 0$, giving $J(u_\infty) = 0 = \min J$ (as $J \geq 0$). By (H2), $J(u^*) = \frac{1}{2} \|r^*\|^2 > 0$. (ii) $\nabla J(u) = DR_h(u)^\top R_h(u)$; evaluate at u^* and use (7). (iii) $\frac{1}{2} \|r^* + Le\|^2$ is convex in e ; its normal equations $L^\top Le = -L^\top r^*$ give $e_\infty = -(L^\top L)^{-1} L^\top r^* = -L^+ r^*$ ($L^\top L$ invertible by (H3)), and $\|L^+\|_2 = 1/\sigma_{\min}(L)$. (iv) To first order $\dot{e}|_0 = -L^\top(r^* + Le)$, so $\frac{d}{dt} \frac{1}{2} \|e\|^2|_0 = e \cdot \dot{e} = -e^\top L^\top(r^* + Le) = -\|Le\|^2 - (Le) \cdot r^*$. This is positive iff $(Le) \cdot r^* < -\|Le\|^2$; taking $Le = -\alpha r^* / \|r^*\|$ (feasible whenever the projection of r^* onto range L is nonzero) the condition becomes $0 < \alpha < \|r^*\|$, an open set. $\square$

Good detector, bad fixer—from one decomposition. Equation (7) splits the monitored residual into a fixed floor r^* and an error-driven term Le . Far from the truth, $\|Le\| \gg \|r^*\|$, so $\|R_h(\hat{u})\| \approx \|Le\|$, and since $\sigma_{\min}(L)\|e\| \leq \|Le\| \leq \sigma_{\max}(L)\|e\|$ the residual norm is a two-sided proxy for the field error (up to $\kappa(L)$)—this is the calibrated *trust signal* (rank correlation up to 0.83, subsection 5.3). Near the truth,

$\|Le\| \lesssim \|r^*\|$, the floor dominates, and the residual decouples from the error: its minimiser is displaced from u^* by [Theorem 1\(iii\)](#) and a residual-reducing step can grow the error by (iv)—this is the failure of the residual as a correction *objective* ([subsection 5.5](#)). Note the divergence in (iv) vanishes when $r^* = 0$: it is driven by the *floor*, not by ill-conditioning, which merely sets the floor’s magnitude in (iii).

Empirical confirmation. On 200/200 real AirFRANS test cases, the monitored residual of the ground-truth field is substantially nonzero (**(H2) holds**): $\|r^*\|$ has mean 0.192 (median 0.133), while the uniform-freestream field gives $\|R_h(u_\infty)\| = 0$ in every case—the objective strictly prefers a physically wrong field to the truth ([results/certificates/residual_floor_realdata.json](#)). A trained backbone’s over-smoothed prediction sits *below* the truth floor in 160/200 cases, consistent with the loop’s drive away from u^* in [Table 5](#).

Assumptions, stated plainly. The theorem is about the *monitored discrete* operator R_h —what the loop minimises—not the continuous PDE; (H2) holds here because R_h omits the no-slip closure *and* the 128^2 grid under-resolves the boundary layer and drops the spatially-varying- ν_t term $\nabla\nu_t \cdot \nabla u$. (H3) makes the floor a genuine displacement. The detector leg is a *regime* statement: its tightness is controlled by $\kappa(L)$, so we claim a positive rank correlation, not an isometry.

Relation to prior work. [Krishnapriyan et al. \(2021\)](#) tie ill-conditioning to the *optimisation difficulty* of PINN losses, and [Wang et al. \(2022\)](#) analyse spectral bias and loss-term imbalance in PINN *training dynamics*; both concern trainability, not a displacement of the residual minimiser from the truth. The forward-vs-backward error gap $\|e\| \leq \|R_h\|/\sigma_{\min}(L)$ is classical ([Trefethen & Bau, 1997](#)), and a-posteriori PINN bounds upper-bound error *by* residual (the benign regime). To our knowledge the *residual-floor* statement—that a self-consistent monitored residual makes the truth a non-minimiser, yielding “good detector, bad fixer” simultaneously—is new for a learned neural-operator CFD corrector, and it is exactly the dissociation our experiments report.

7 Limitations

- **Backbone-dependent correction quality.** The supervised DEQ corrector reduces field MSE by 9–25% on the SOTA Transolver backbone ([subsection 5.2](#)) but is flat-to-unhelpful on the weak grid backbone ([subsection 5.3](#)). We report both regimes; we do not claim the corrector helps on an arbitrary backbone, only that it helps on a strong one. The trust signal, by contrast, is strong on both ($\rho = 0.625$ SOTA, ≈ 0.40 – 0.83 weak). We make no competitiveness claim for the grid backbone (it trails SOTA ~ 4 – $60\times$); the Transolver backbone we run *on* is itself competitive.
- **Grid resolution.** A uniform 128^2 Cartesian grid cannot resolve a $\text{Re} \approx 10^6$ boundary layer (sub-cell), so wall quantities are approximate and the absolute MSE values are not comparable to body-fitted solvers.
- **Single dataset / 2-D.** All results are on 2-D AirFRANS airfoils. 2-D bluff bodies and 3-D geometries (AhmedML, [Ashton et al., 2024](#); DrivAerNet++, [Elrefaie et al., 2024](#)) are roadmap targets; the phenomenon has not been tested beyond AirFRANS.
- **Conformal coverage out-of-distribution is empirical, not guaranteed.** The distribution-free guarantee assumes exchangeability, which fails under shift; OOD coverage holds in our experiments only because the uncertainty inflates appropriately ([subsection 5.6](#)). Scores are also pooled over spatially-correlated cells, so the effective sample size is below the raw cell count, and the calibration set is small ($n = 100$); a larger-calibration-set robustness check is future work.
- **Residual-as-objective remains a dead end.** Used as a correction objective or acceptance gate (not as a supervised input feature), the residual still fails: more iterations raise the PDE residual while error falls, and the monotone-residual acceptance test accepts almost no steps ([subsection 5.5](#)). The DEQ contraction guarantee is in the correction variable δ , and the acceptance test’s monotone-residual guarantee is real but vacuous as an accuracy mechanism. The self-correction gains in [subsection 5.2](#) come from the *supervised* input-feature path, not from minimising the residual.

- **Adaptive uncertainty on the SOTA model requires a deep ensemble.** On the near-deterministic Transolver backbone, MC-dropout σ at dropout 0.05 is near-degenerate, so although conformal coverage holds (0.902, target 0.90), the quantile q is enormous ($\approx 10^9$) and ECE ≈ 0.31 : coverage there is delivered by a near-constant band. A deep ensemble of $M=5$ members recovers a genuinely input-adaptive band on this model ($q = 2.35$, ECE 0.074; Table 2), at $5\times$ the training cost. That ensemble certificate is reported as a single run (one cal/test split), not a 3-seed average; multi-seed error bars on the ensemble ECE and a study of M are future work.
- **Eddy-viscosity channel near-degenerate.** The ν_t output is effectively unlearned at this scale (per-arm MSE $\approx 5 \times 10^{-8}$); since residuals use $\nu_{\text{eff}} = \nu + \nu_t$, the physics term currently leans on the laminar viscosity. A trainable ν_t target is future work.
- **Classical fallback is a stub interface.** The trust-gated fallback fixes the integration seam and reports what would run; the OpenFOAM/SU2 backends are not implemented and bear on no quantitative claim.
- **Sensitivity sweeps are light.** The iteration and toggle sweeps (subsection 5.5) are single-checkpoint, unseeded sweeps; they illustrate the residual-error divergence robustly in direction but are not multi-seed.

8 Conclusion

We set out to use the steady-RANS physics residual of a neural-CFD surrogate, and we found that it plays **three distinct roles with three distinct verdicts**. *As a trust signal* it is an excellent, backbone-agnostic detector of error: $\rho = 0.625 \pm 0.019$ on a SOTA Transolver backbone, $\approx 0.40\text{--}0.83$ on a weak backbone, and—under distribution shift on the weak backbone—regime-invariant (≈ 0.74 , rescuing an **aoa** collapse from 0.31 to 0.75) where the one-shot signal collapses. *As an input feature to a supervised corrector* it works on a strong backbone: the DEQ corrector cuts volume field MSE by 9–25% on all 3 seeds of the SOTA Transolver model, with force ranking held. *As a correction objective or acceptance gate* it fails: more iterations raise the residual while error falls, and the monotone-residual gate accepts almost nothing—minimising the residual is not minimising the error. On the strength of the detector result we package a backbone-agnostic conformal trust layer that holds target coverage both in-distribution (0.91/0.93/0.94 at 0.90, weak backbone) and on the SOTA backbone (0.902), and empirically retains coverage out-of-distribution via uncertainty inflation; pairing it with a deep-ensemble σ makes the band input-adaptive even on the near-deterministic SOTA model ($q = 2.35$, ECE 0.074), where MC-dropout alone gives a near-constant one. One honest caveat carries the tension: correction quality is backbone-dependent (flat on the weak backbone). The durable contribution is the self-auditing, backbone-agnostic trust layer and the supervised corrector it accompanies—demonstrated on a competitive surrogate, reproducible end-to-end from the committed harness.

References

- Anonymous. Learned residual error correction for numerical pde surrogates. *arXiv preprint arXiv:2306.12047*, 2023.
- Neil Ashton, Danielle Maddix, Samuel Gundry, and Parisa Shabestari. AhmedML: High-fidelity computational fluid dynamics dataset for incompressible, low-speed turbulent flows. *arXiv preprint arXiv:2407.20801*, 2024.
- Shaojie Bai, J. Zico Kolter, and Vladlen Koltun. Deep equilibrium models. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2019. arXiv:1909.01377.
- Florent Bonnet, Jocelyn Ahmed Mazari, Paola Cinnella, and Patrick Gallinari. AirfRANS: High fidelity computational fluid dynamics dataset for approximating reynolds-averaged navier-stokes solutions. In *Advances in Neural Information Processing Systems (NeurIPS) Datasets and Benchmarks Track*, 2022. arXiv:2212.07564.

-
- Mohamed Elrefaie, Florin Morar, Angela Dai, and Faez Ahmed. DrivAerNet++: A large-scale multimodal car dataset with computational fluid dynamics simulations and deep learning benchmarks. *arXiv preprint arXiv:2406.09624*, 2024.
- Samy Wu Fung, Howard Heaton, Qiuwei Li, Daniel McKenzie, Stanley Osher, and Wotao Yin. JFB: Jacobian-free backpropagation for implicit networks. In *AAAI Conference on Artificial Intelligence*, 2022. arXiv:2103.12803.
- Yarin Gal and Zoubin Ghahramani. Dropout as a bayesian approximation: Representing model uncertainty in deep learning. In *International Conference on Machine Learning (ICML)*, 2016. arXiv:1506.02142.
- Jun-Ting Hsieh, Shengjia Zhao, Stephan Eismann, Lucia Mirabella, and Stefano Ermon. Learning neural PDE solvers with convergence guarantees. In *International Conference on Learning Representations (ICLR)*, 2019. arXiv:1906.01200.
- Aditi S. Krishnapriyan, Amir Gholami, Shandian Zhe, Robert M. Kirby, and Michael W. Mahoney. Characterizing possible failure modes in physics-informed neural networks. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2021.
- Balaji Lakshminarayanan, Alexander Pritzel, and Charles Blundell. Simple and scalable predictive uncertainty estimation using deep ensembles. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2017. arXiv:1612.01474.
- Zongyi Li, Nikola Kovachki, Kamyar Azizzadenesheli, Burigede Liu, Kaushik Bhattacharya, Andrew Stuart, and Anima Anandkumar. Fourier neural operator for parametric partial differential equations. In *International Conference on Learning Representations (ICLR)*, 2021. arXiv:2010.08895.
- Zongyi Li, Daniel Z. Huang, Burigede Liu, and Anima Anandkumar. Fourier neural operator with learned deformations for pdes on general geometries. *arXiv preprint arXiv:2207.05209*, 2022.
- Phillip Lippe, Bastiaan S. Veeling, Paris Perdikaris, Richard E. Turner, and Johannes Brandstetter. PDE-refiner: Achieving accurate long rollouts with neural PDE solvers. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023. arXiv:2308.05732.
- Huakun Luo, Haixu Wu, Hang Zhou, Lanxiang Xing, Yichen Di, Jianmin Wang, and Mingsheng Long. Transolver++: An accurate neural solver for PDEs on million-scale geometries. *arXiv preprint arXiv:2502.02414*, 2025.
- Ziqi Ma, Kamyar Azizzadenesheli, and Anima Anandkumar. Calibrated uncertainty quantification for operator learning via conformal prediction (UQNO). In *International Conference on Learning Representations (ICLR)*, 2024. arXiv:2402.01960.
- Tanya Marwah, Ashwini Pokle, J. Zico Kolter, Zachary C. Lipton, Jianfeng Lu, and Andrej Risteski. Deep equilibrium based neural operators for steady-state PDEs. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023. arXiv:2312.00234.
- Maziar Raissi, Paris Perdikaris, and George E. Karniadakis. Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving nonlinear partial differential equations. *Journal of Computational Physics*, 378:686–707, 2019.
- Rishikesh Ranade, Mohammad Amin Nabian, Kaustubh Tangsali, Ali Kashefi, Max Rietmann, Sanghyun Kim, and Sanjay Choudhry. DoMINO: A decomposable multi-scale iterative neural operator for external aerodynamics. *arXiv preprint arXiv:2501.13350*, 2025. NVIDIA.
- Lloyd N. Trefethen and David Bau. *Numerical Linear Algebra*. SIAM, 1997.
- Sifan Wang, Xinling Yu, and Paris Perdikaris. When and why PINNs fail to train: A neural tangent kernel perspective. *Journal of Computational Physics*, 449:110768, 2022.

Ezra Winston and J. Zico Kolter. Monotone operator equilibrium networks. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2020. arXiv:2006.08591.

Haixu Wu, Huakun Luo, Haowen Wang, Jianmin Wang, and Mingsheng Long. Transolver: A fast transformer solver for PDEs on general geometries. In *International Conference on Machine Learning (ICML)*, 2024. arXiv:2402.02366.